

Lecture 7

COMP 3760

Data Structures and Graphs

Text chapter 1.4, 3.5, 5.3

Fundamental Data Structures

(Chapter 1.4)

Data Structures

- A *data structure* is a particular way of storing and organizing data
- Data structures and algorithms are often deeply interconnected
 - The way you organize data affects the performance of your algorithm
- We've *mostly* been using arrays ... so far

Fundamental Data Structures

- Linear Data Structures
 - Array
 - Linked list
 - Stack
 - Queue
- Set
- Dictionary (Map)
- Tree
- Graph

Arrays

- A sequence of n items of the same type, accessed by an index




- The good:
 - Each item accessed in same constant time
- The bad:
 - Size is fixed
 - Insertion / deletion in an array is time consuming – all the elements following the inserted element must be shifted appropriately

Linked Lists

- (singly) A sequence of zero or more elements called *nodes*, consisting of data and a pointer




- (doubly) Pointers in each direction



Linked Lists


- Linked lists provide two key advantages over arrays
 - Dynamic size
 - Ease of insertion/deletion
- Linked lists have some drawbacks:
 - Random access is not allowed

Do you know what this is?



Stack

- Like a stack of plates
- Last-in-first-out (LIFO)



Operations on a stack

- Insert operation is called Push
- Delete operation is called Pop
- Examining the top item is Peek
- Example application:
 - Analysis of languages (e.g. properly nested brackets)
 - Properly nested: `(())`
 - Wrongly nested: `(())(())`

Stack

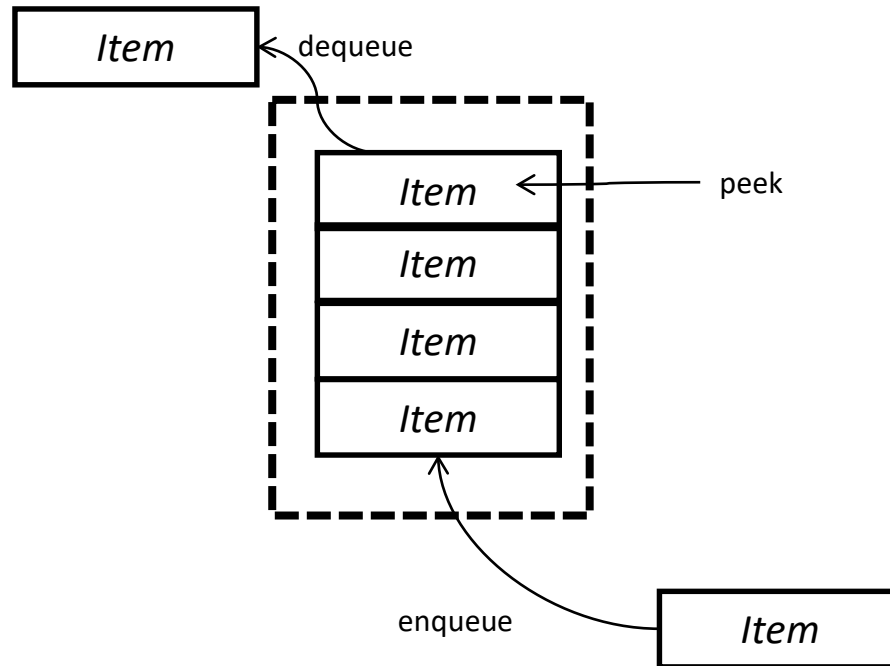
Algorithm CheckBalancedParenthesis(expr)

```
1.  n = length(expr)
2.  Create a stack s
3.  for i = 0 to n-1 do
4.      if (expr[i] is '(' ) do
5.          s.Push(expr[i])
6.      else if (expr[i] is ')' )
7.          if (s is empty) or
              (s.Peek() does not pair with expr[i])
8.              return False
9.          else
10.             s.Pop()
11. if (s is empty)
12.     return True
13. else
14.     return False
```

Abstract Data Type

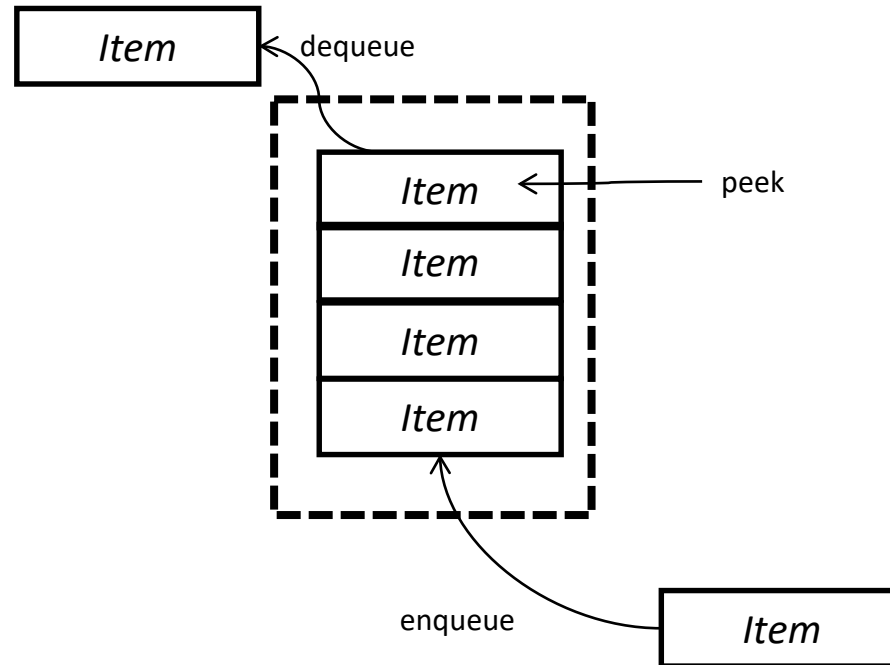
- Often a data structure is closely associated with a set of available operations
- Data structure + operations = *abstract data type*
 - From an OOP perspective, think about members (methods) of a class
- Example 1: priority queue
 - Underlying implementation was a heap
 - Operations were Insert and deleteMax
- Example 2: stack
 - Operations are push, pop, peek

How about this one?



Queues

- Like a line-up
- First-in-first-out (FIFO)



Operations on a queue

- Adding to the queue is Enqueue
- Removing from the queue is Dequeue
- The top/front element is the Head (sometimes there is a “Peek” method)

Set

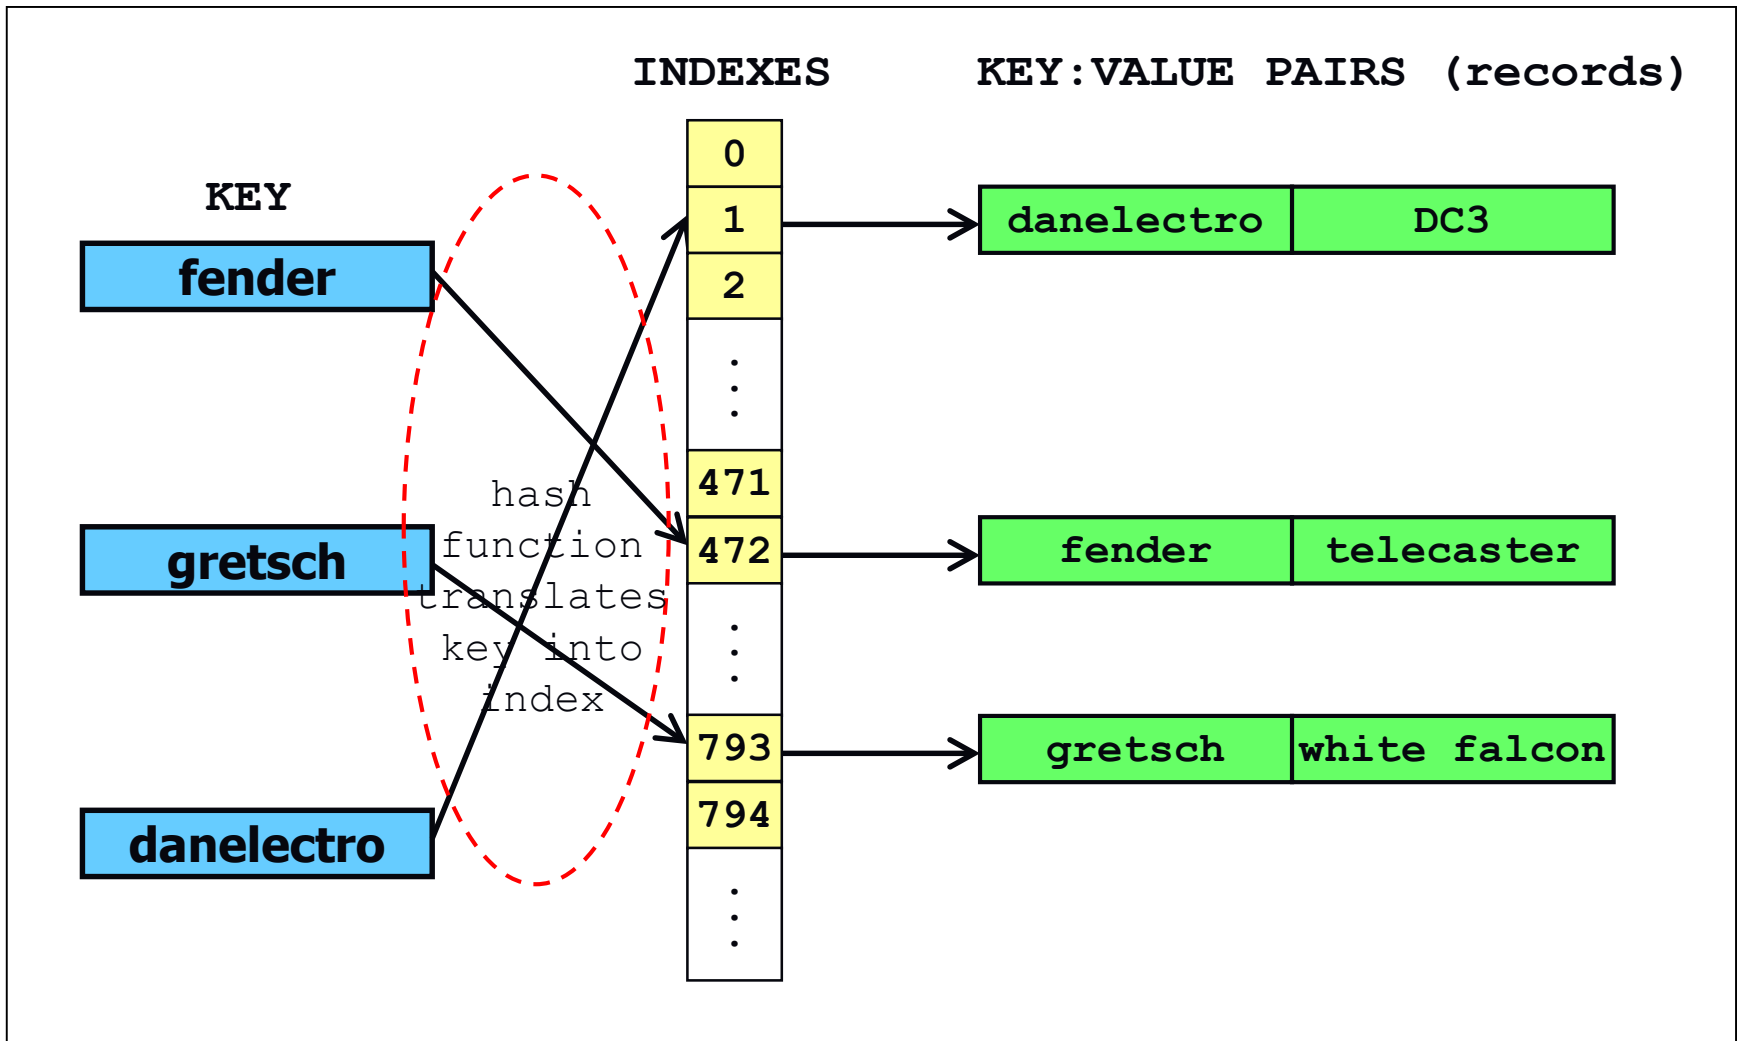
- Like a set in math, e.g. {1, 2, 3, 4}
 - Sets cannot contain duplicate items
- Operations on a Set:
 - Add an item to it
 - Remove an item from it
 - Check if an item is in it
 - Iterate over it (loop based on all items, one-by-one)

Set in Java

- Different ways to implement a set
 - HashSet
 - Faster implementation, but it is unordered
 - TreeSet
 - Slower, but the items are available in order

Map (as a hash table)

- A **Map** is a lookup table that takes a **key** and returns a **value**
 - the most common implementation is as a hashtable (hashmap)




Graphs


(Still in Chapter 1.4)

Graphs


- $G = (V, E)$
 - V is a set of *vertices*
 - E is a set of *edges*



Motivation: Real world connections



Undirected




Directed

Some special graphs

- *Connected graph*
 - A graph where there is a path available between any two vertices
- *Cyclic graph*
 - A graph containing at least one cycle
- *Acyclic graph*
 - A graph containing no cycles
- *Tree*
 - Any connected + acyclic graph
- *Complete graph*
 - Every pair of vertices is connected by an edge
- *Weighted graph*
 - Every edge has an associated value

Trees

- Connected, acyclic graphs
 - Usually we think of trees as having a root
- Representing data in a tree can speed up your algorithms in many natural problems




Representation of graphs

- Two common ways to represent graphs:
 - Adjacency matrix
 - $|V| \times |V|$ matrix
 - Cell i, j represents an edge from vertex i to j
 - Adjacency lists
 - $|V|$ linked lists – one for each vertex, showing all the neighbours of that vertex

Representation: Adjacency Matrix

- ▶ For this graph:




- Adjacency matrix is the following:

$O(|V|^2)$

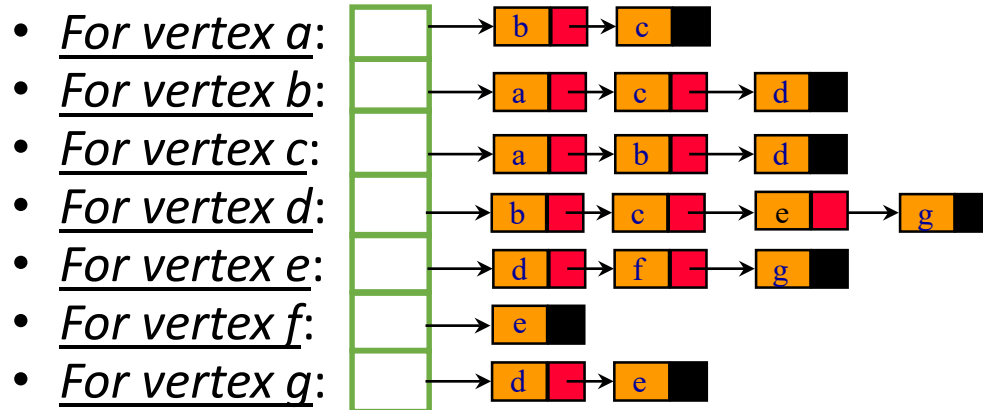
	a	b	c	d	e	f	g
a	0	1	1	0	0	0	0
b	1	0	1	1	0	0	0
c	1	1	0	1	0	0	0
d	0	1	1	0	1	0	1
e	0	0	0	1	0	1	1
f	0	0	0	0	1	0	0
g	0	0	0	1	1	0	0

Representation: Adjacency List

- ▶ For the same graph:



- Adjacency list is the following:



$O(|E|)$

Representing Graphs

1. Adjacency matrix

- Or Weight Matrix for weighted graphs

2. Adjacency lists


- A list of vertices connected to each vertex
- Which one to use?
 - Depends on the nature of the graph (sparse or not)
 - Depends on the algorithm

Graph Algorithms

(Chapter 3.5)

Graph Traversal

- Many real-world problems require processing of each vertex (or edge) in a graph



- Routing a message on a network
- Web crawling
- Social networking
- Garbage collection
- Solving puzzles

Graph Traversal Algorithms


- Graph traversal algorithms give a method for *systematically processing* all vertices

Basic idea: "visit" all the vertices, one at a time, *marking* them as we visit them

- Two approaches:
 - Depth-First Search (DFS)
 - Breadth-First Search (BFS)

Depth-first search (DFS)


- Think about how you might try to find your way through a maze



Depth-first search (DFS)

- Visits all vertices by always moving away from the last vertex visited (if possible)
 - Backtracks at “dead-ends” (no adjacent, unvisited vertices)
- Implementation often uses a stack of vertices being processed
- Follows a tree-like route throughout the graph

DFS example



Backtrack/finish order: e f h d c g b a

DFS order: a b f e g c d h

Some notes on DFS

- To track the progress of the algorithm we use a stack
 - When we make a recursive call, e.g. $\text{dfs}(v)$, we push v onto the stack
 - When v is a dead-end (i.e. no more neighbors to visit) it is popped off the stack
- Our convention: break ties for “next neighbor” by using some natural order
- Typical results from running DFS can be:
 - List of vertices in order visited
 - List of vertices in order of “dead-ends” (when popped from stack)
 - DFS Tree – tree containing all the edges that were used to visit nodes
 - Unused edges of G (edges not in DFS tree) are called “back edges”

DFS algorithm

```
Algorithm Depth_First_Search(Graph G)
// Graph G = {V,E}
  initialize visited to false for all vertices
  for each vertex v in V
    if v has not been visited
      dfs_helper(v)

function dfs_helper(Vertex v)
  visit node v
  for each vertex w in V adjacent to v
    if w has not been visited
      dfs_helper(w)
```

- “Visit node v” means doing whatever you need to do at each node

Common uses of DFS

- Find a spanning tree of a graph
- Find a path between two vertices v and u
- Find a path out of a maze
- Determine if a graph has a cycle
- Find all the connected components of a graph
- Search the state-space of problems for a solution (AI)
- Many more!

Efficiency of DFS

- The basic operation is the if statement in dfs_helper():

```
for each vertex w in V adjacent to v
    if w has not been visited
        dfs(w)
```

- Each time dfs_helper() is called, this loop examines all the neighbors of some vertex v ... eventually it looks at ALL the neighbors of ALL the vertices
 - Therefore the number of basic operations depends on the data structure used to implement the graph
- Basically we need to visit each element of the data structure exactly once. So the efficiency must be:
 - $O(|V|^2)$ for adjacency matrix
 - $O(|V| + |E|)$ for adjacency lists


Which is better/worse?

- $O(|V|^2)$
- $O(|V| + |E|)$


Breadth-first search (BFS)


- Recall the “trick” where we used an array to store a (very specific type of) binary tree (a heap).

Breadth First Search
is sorta like this
—“level by level” —
but it’s about
*TRAVERSING
A GRAPH*



17 11 12 9 8 10 5 1 4 6 2 3 7






Breadth-first search (BFS)

- Visit all neighbors “the same distance” from starting vertex
 - Visit immediate neighbors first
 - Then the neighbors of all those vertices
 - Etc.
- Instead of a stack, BFS uses a queue
- Follows a tree-like route throughout the graph, but perhaps a different tree than DFS

BFS example



BFS order: a b e f g c h d

BFS algorithm

```
Algorithm Breadth_First_Search(Graph G)
// Graph G = {V,E}
  initialize visited to false for all vertices
  for each vertex v in V
    if v has not been visited
      bfs_helper(v)

function bfs_helper(Vertex v)
  visit node v
  initialize a queue Q
  add v to Q
  while Q is not empty
    for each w adjacent to Q.head
      if w has not been visited
        visit node w
        add w to Q
  Q.dequeue()
```

- Uses a queue (FIFO) to determine which vertex to visit next
- Edges that are in G but not in the resulting BFS tree are called “cross-edges”

Notes on BFS

- Same efficiency as DFS:
 - Adjacency matrix: $O(|V|^2)$
 - Adjacency list: $O(|V| + |E|)$
- Yields just one ordering of vertices (order added/deleted from queue is the same)
 - Whereas with DFS, the order that vertices are *visited* may be different from the order they get *finished* (become dead-ends)

BFS applications

- Really the same as DFS
- Sometimes one or the other may be better for specific problems

Graph Algorithms: Binary tree traversal


Textbook: Chapter 5.3

Tree traversal

- Traversing a tree means to visit all of the nodes of the tree
- We've already seen DFS and BFS (for graphs)
- Here are a few traversals specific to *binary trees*:
 - Preorder – root *before* the children
 - Inorder – root *between* the children
 - Postorder – root *after* the children

Preorder traversal


1. Visit the root
2. Traverse the left subtree
3. Traverse the right subtree



Preorder traversal is: a b c

Inorder traversal


1. Traverse the left subtree
2. Visit the root
3. Traverse the right subtree



Inorder traversal is: b a c

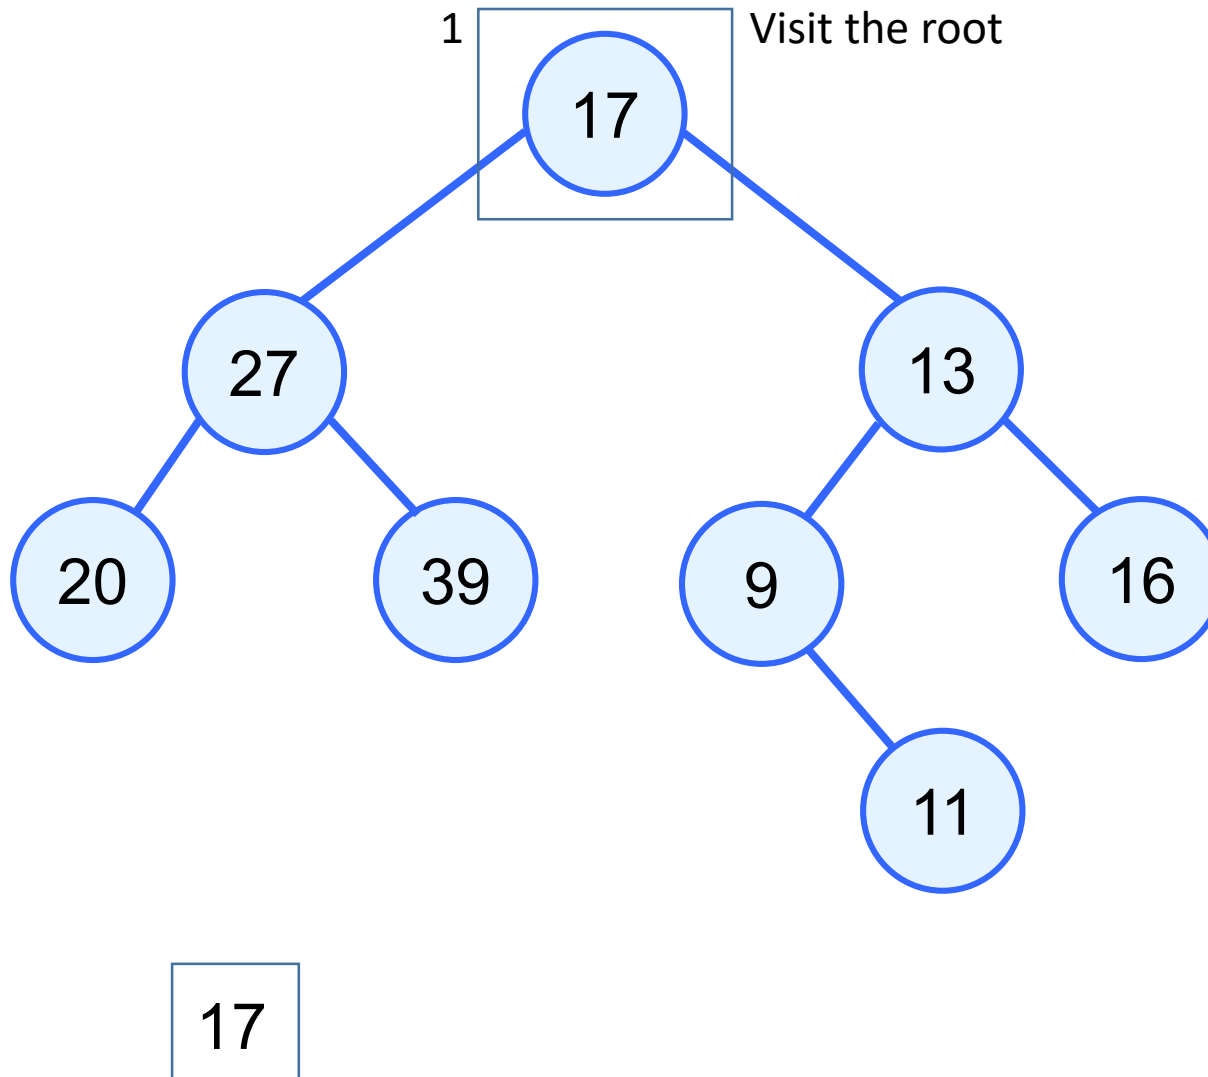
Postorder traversal

1. Traverse the left subtree
2. Traverse the right subtree
3. Visit the root

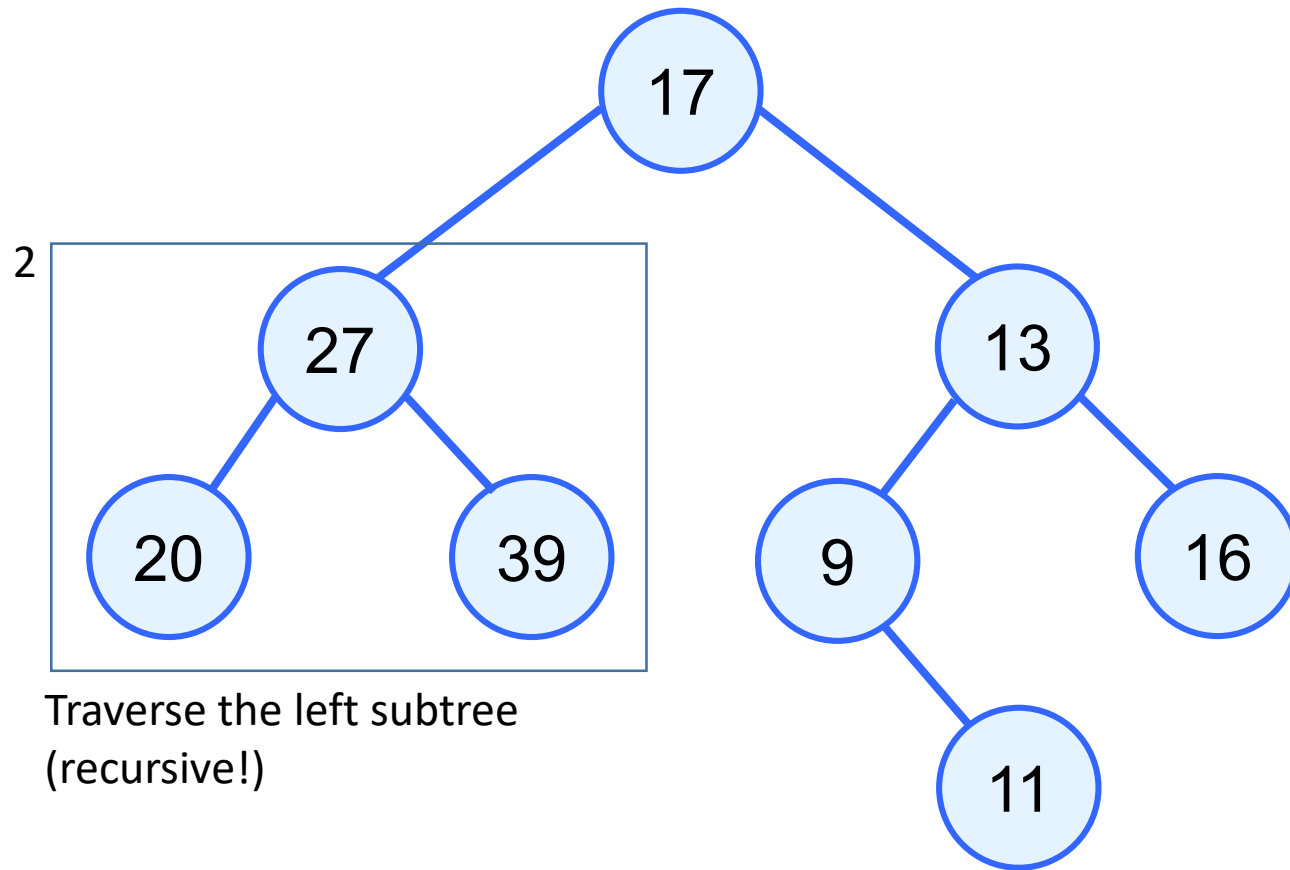


Postorder traversal is: b c a

Preorder example



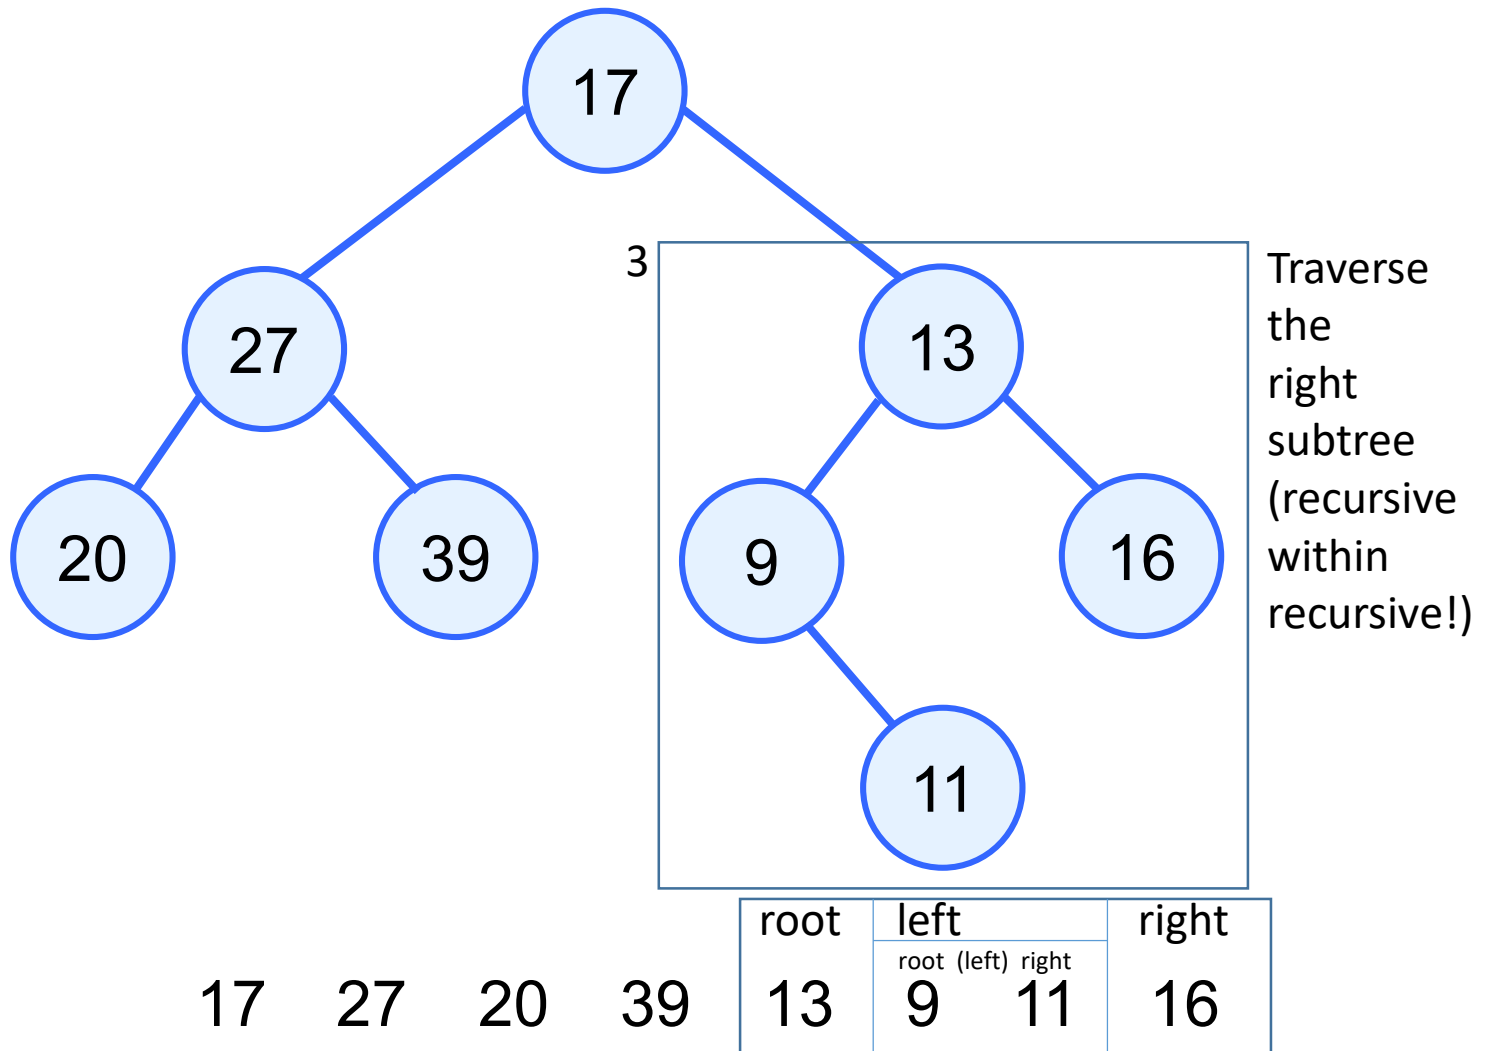
Preorder example




Traverse the left subtree
(recursive!)

17	root	left	right
	27	20	39

Preorder example




Preorder example



17 27 20 39 13 9 11 16


Another way to think about it

- Consider this path that meanders past all of the nodes
- ALL three traversals follow this path!




Preorder traversal

- Is when the path passes the LEFT side of each node



Inorder traversal



- Is when the path passes the **BOTTOM** of each node



Inorder: D B E A F C G

Postorder traversal

- Is when the path passes the **RIGHT** side of each node



Postorder: D E B F G C A

Pseudocode

```
Algorithm inOrder (Node N)
if N != null
    inOrder(N.leftChild)
    Print N.value
    inOrder(N.rightChild)
```

```
Algorithm preOrder (Node N)
if N != null
    Print N.value
    preOrder(N.leftChild)
    preOrder(N.rightChild)
```

```
Algorithm postOrder (Node N)
if N != null
    postOrder(N.leftChild)
    postOrder(N.rightChild)
    Print N.value
```

What if I told you

Preorder → A B D E C F G

Inorder → D B E A F C G

Fun facts about pre/in/postorder

- Given pre + in, you can reconstruct the tree
 - (and also determine postorder)
- Given post + in, you can reconstruct the tree
 - (and also determine preorder)
- Given pre + post, you can only *sometimes* reconstruct the tree
 - For you to ponder: under what condition(s)?


Some problems solvable
with graph traversal

Problem 1: Spanning Tree

- Given a connected graph G , find a spanning tree T
 - This is a straight-up application of BFS (or DFS)
 - Build a new graph (the spanning tree) as we go
 - Initialize a new graph T
 - Each time we visit a vertex, add the edge we used to T
- BFS or DFS?
 - BFS usually gives shorter paths between vertices

Problem 2: Solving a Maze

- Represent maze as a graph
 - Nodes for start, finish, intersections, and dead-ends
 - Find a path from *start* to *finish*
- BFS or DFS?
 - If interested in end-result:
 - BFS will find the shortest total path
 - If actually walking while solving:
 - DFS tends to result in less actual walking
 - BFS backtracks to parent nodes too often



Problem 3: Shortest Path

- Find the shortest path between two vertices u and w
- BFS or DFS?
 - BFS will find a shortest path
 - DFS will find a path – maybe not the shortest one
- Idea of algorithm (and why it works):
 - First, use $\text{bfs}(u)$ to create a spanning tree T with u as the root. Note that all paths that appear in T are the shortest paths from u to their respective vertices
 - Then, use DFS on T to find a path from u to w (as in the maze problem)

Problem 4: Determine Connectivity

- Determine if a graph is connected
- BFS or DFS?
 - Either will work
 - Think about this yourself!
 - What modification(s) do you need to make to the algorithm to answer this question?