

# Lecture 9

COMP 3760

Greedy Algorithms (Text chapter 9)

# On tap for today

- Random bits
  - Graph algorithms *meta* recap
- Greedy algorithms
  - Idea of greedy algorithms
  - Ex: Minimum spanning trees (MST)
    - Prim's algorithm
    - Kruskal's algorithm
  - Ex: Single-source shortest paths (SSSP)
    - Dijkstra's algorithm

## Solving problems using graphs

- Strategy 1
  - Modify existing graph algorithm
- Strategy 2
  - Represent the problem as a “clever” graph
  - Feed the graph to a Graph Algorithm
  - Use the output to determine the answer to your problem
- Common reality:
  - Hybrid of those two

# Known graph algorithms (so far)

- DFS of  $G$ 
  - Output options: DFS, dead-end, spanning tree
- BFS of  $G$ 
  - Output options: BFS order, spanning tree
- Count connected components of  $G$ 
  - Modified DFS/BFS
- Is  $G$  connected?
  - Modified DFS/BFS
- Topological sort of  $G$ 
  - 1) Modified DFS; 2) Custom algorithm (dec&conq)
- New algorithms today (greedy ones):
  - Minimum spanning tree (MST) of  $G$  (Prim, Kruskal)
  - Single-source shortest paths (SSSP) of  $G$  (Dijkstra)

## Problems we've solved

- Counting regions on a map of coloured hex tiles
  - How: Counting connected components
- Getting dressed in the morning
  - How: Topological sorting
- Connecting a set of locations all together via a physical network
  - How: MST


Onward with  
greedy algorithms

# Making change

- Imagine you're a shop clerk giving change and you want to use the *smallest number of coins*
- Strategy:
  - Always select the biggest feasible coin
- Example: 37 cents
  - 1 quarter (need 12 more cents)
  - 1 dime (need 2 more cents)
  - 2 pennies (2 *what?*)



This is a “greedy algorithm”

A photograph of several children sitting on a light-colored carpet, playing with colorful plastic toys. The toys include a green bell-shaped object, a red and yellow toy, and a red toy with a black handle. The children are wearing casual clothing, and their legs and feet are visible. The image is slightly faded, serving as a background for the text.

Always make the choice  
that looks best *right now*




# Making change—algorithm

```
Algorithm MakeChange(N)
    sum = 0
    coins = {}      // set of coins to be returned
    while sum < N do
        choose the largest coin X with value <= (N-sum)
        sum += X.value
        coins += {X}
    endwhile
    return coins
END
```

# Does this algorithm always give the best result?

- For US/Canadian coins, yes
  - With or without pennies


- But what if your coins were:



- And you had to give 28 cents?
  - Greedy algorithm result: 25, 1, 1, 1  $\rightarrow$  4 coins
  - But there is a 3-coin answer

# Puzzle

- Make a “smaller” counterexample
  - What if your coins were:



- And you had to make    x    cents?
- I.e., find  $a, b, c, x$  so that the greedy algorithm gives a 3-coin answer, even though a 2-coin answer exists



# Moral of the story

- Greedy algorithms do not always give optimal general solutions to problems
- But sometimes they do
  - and they are “easy” and (if done right) efficient

# Optimization problems and decision problems

- An **optimization problem** is one in which you want to find not just *any* solution, but the *best* solution
  - As opposed to **decision problem** – “does a solution exist?”
  - Decision problem has a yes/no answer
  - Optimization problem is about minimizing or maximizing
- Greedy algorithms attempt to solve *optimization problems*

# Remember the Knapsack problem

- Optimization version:
  - Given  $N$  items with weights + values, and a knapsack with carrying capacity  $W$ , what is the greatest overall *value* of stuff the thief can steal?
- Decision version:
  - Given  $N$  items with weights + values, and a knapsack with carrying capacity  $W$ , can the thief steal  $\$V$  worth of stuff?

# Greedy algorithms

- For solving *optimization problems*
- Construct a solution through a sequence of choices
- Always choose the best option available “right now”
  - The “best” choice is the one that gets us closest to an optimal solution (e.g. take the biggest feasible coin)
- You hope that by choosing a *local* optimum at each step, you will end up at a *global* optimum

# Greedy algorithms

- Greedy choice properties:
  - *Feasible*: Must satisfy the problem's constraints
    - If you are making change for 17 cents, you don't pick a quarter
  - *Locally optimal*: Best local choice among all feasible choices available on that step
    - If you are making change for 14 cents, you pick a dime, not a nickel
    - Assumption: it is “reasonably efficient” to determine this (think about the Knapsack Problem – how to find the “best” choice)
  - *Irrevocable*: Once made, it cannot be changed during subsequent steps of the algorithm



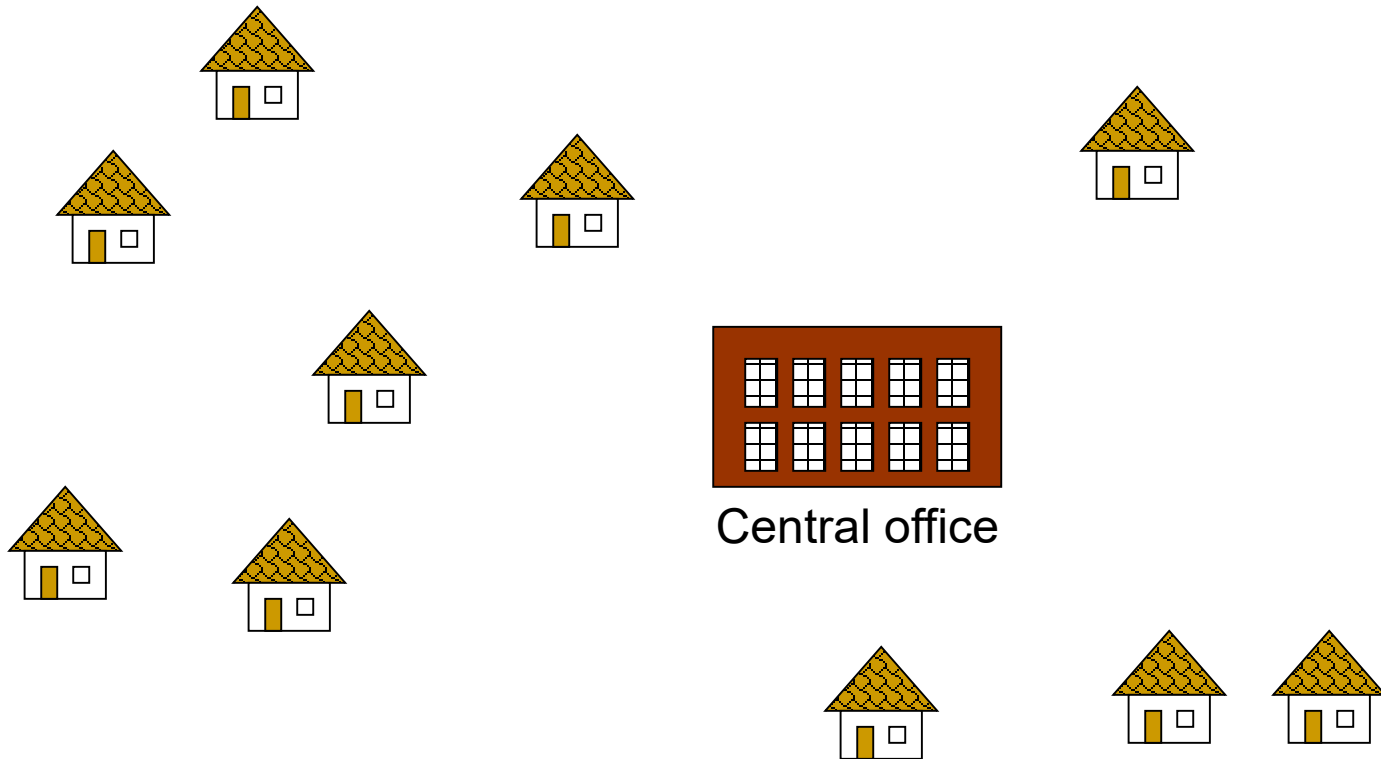
# Greedy algorithms

- We will examine greedy algorithms for the following problems:
  - Finding a minimum spanning tree (MST) of a graph
    - Prim's algorithm
    - Kruskal's algorithm
  - Finding Shortest Paths from a Single Source in a graph
    - Dijkstra's algorithm
  - Coloring a graph

# Greedy algorithm TL/DR

1. Iteratively construct a solution
2. At each step select the “best” item to add
  - Idea for how to select the best should be “simple”

# A real-world problem: Build a (physical) network




# Minimum Spanning Trees


- A **minimum spanning tree** (MST) is a subgraph of a connected, undirected, weighted graph  $G$ , such that
  - it includes all the vertices (“**spanning**”)
  - it is acyclic (“**tree**”)
  - the total cost associated with the edges is the **minimum** among all possible spanning trees
- MST may not be unique

## MSTs (cont'd)

Consider all the  
spanning trees of G:



The weight of each spanning  
tree is given by the sum of its  
edges ...



*Minimum Spanning Tree* of G is this  
graph, and it has a weight of 4.


# If you find MST of a complete graph:

- The result:
  - Is a *tree* (obvs)
    - therefore *connected*
  - connects *all the nodes*
  - using *the minimum cost*

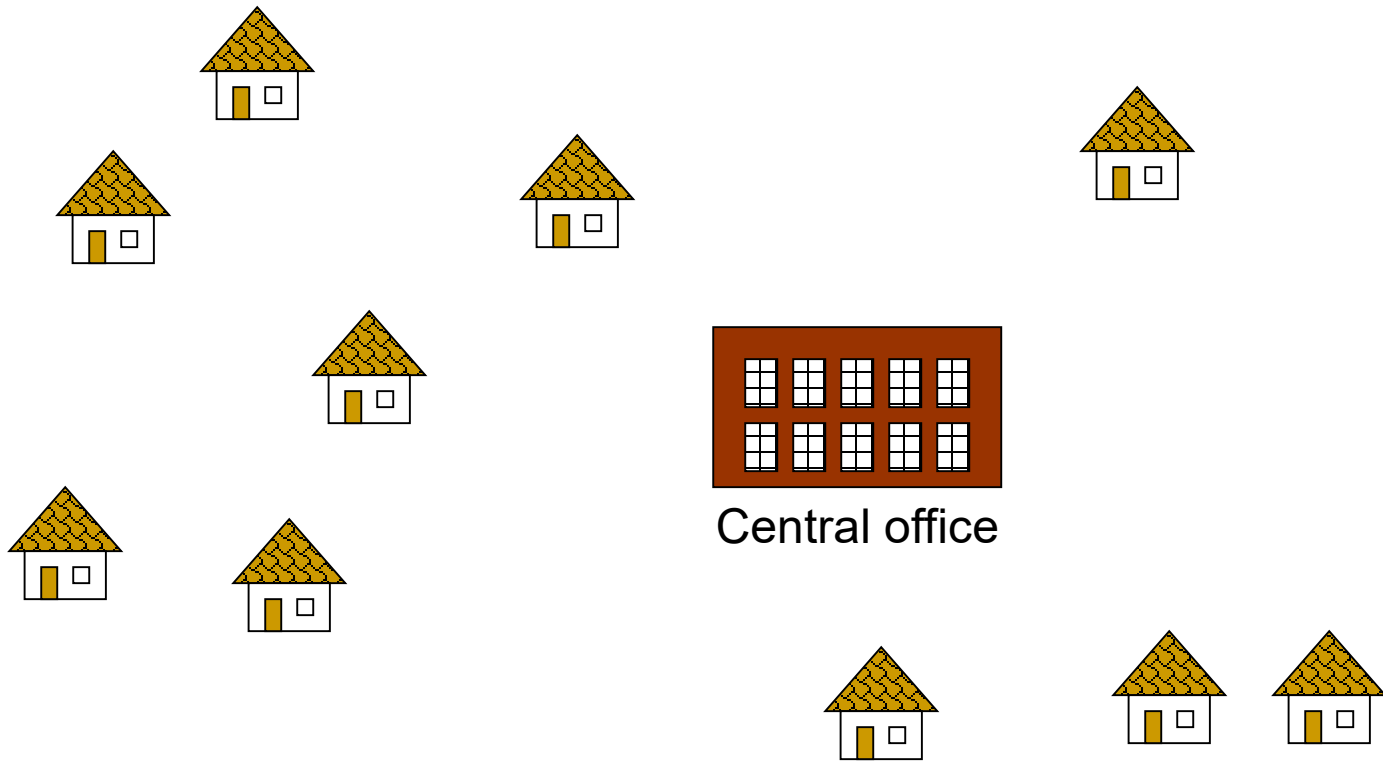
# Reminder: Solving problems with graphs, strategy 2

1. Represent the problem “cleverly” as a graph
2. Feed the graph to a Graph Algorithm
3. Use the output to determine the answer to your problem

We also used Strategy 2 with the “counting map regions” problem (different Graph Algorithm)



# Back to our little village

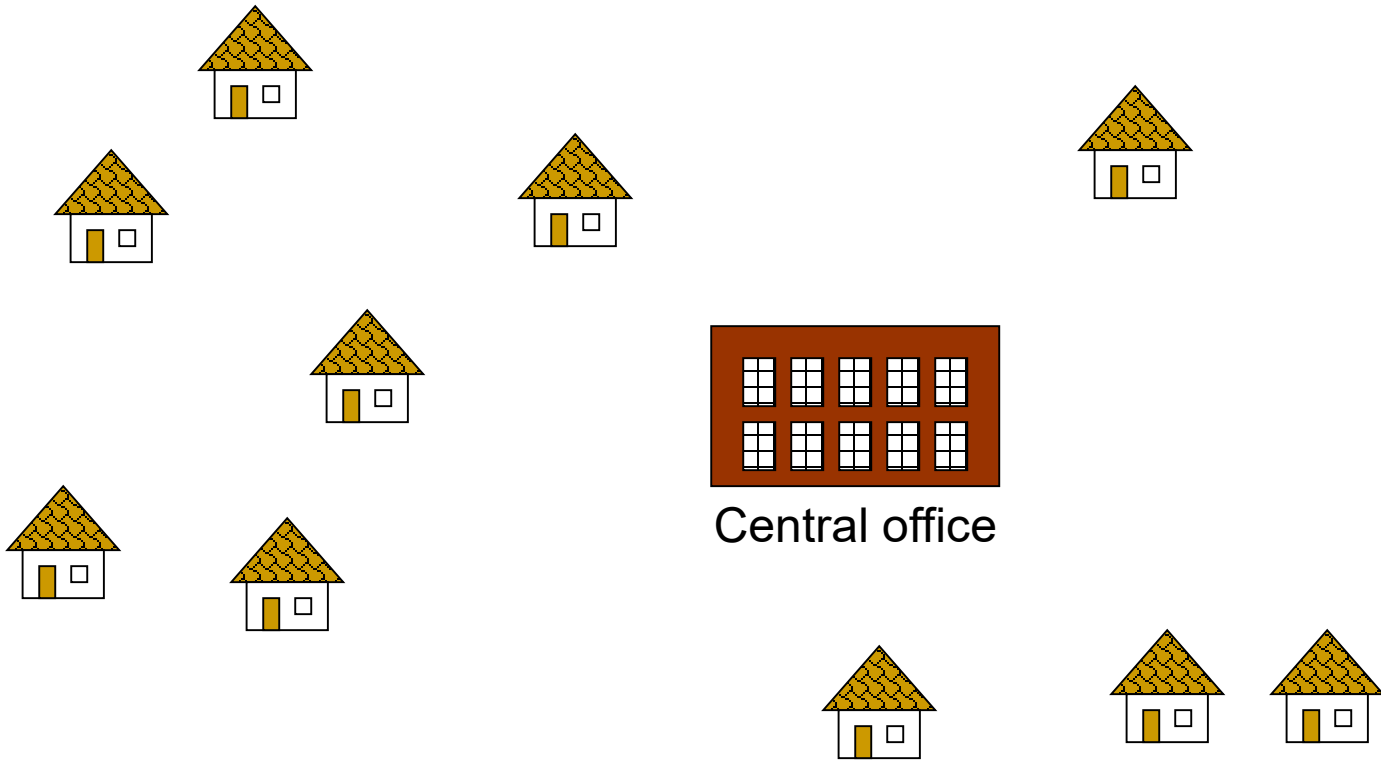



## Let's solve this problem using MST




# Represent it as a graph

- Vertices are all the nodes to be connected
- One edge for *every possible* connection
  - I.e. the complete graph of N vertices
- Each edge has a “weight” associated with it
  - Cost of running a wire from node A to node B
- Now find the MST
  - How does this solve the problem?
    - Spanning tree → all nodes are connected
    - Lowest cost tree → cheapest possible network



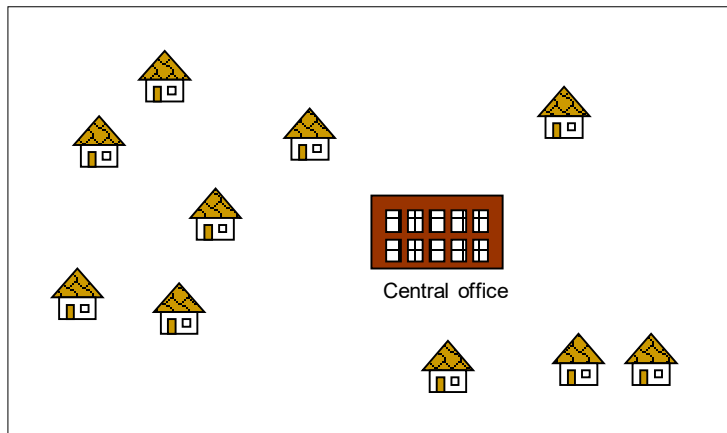






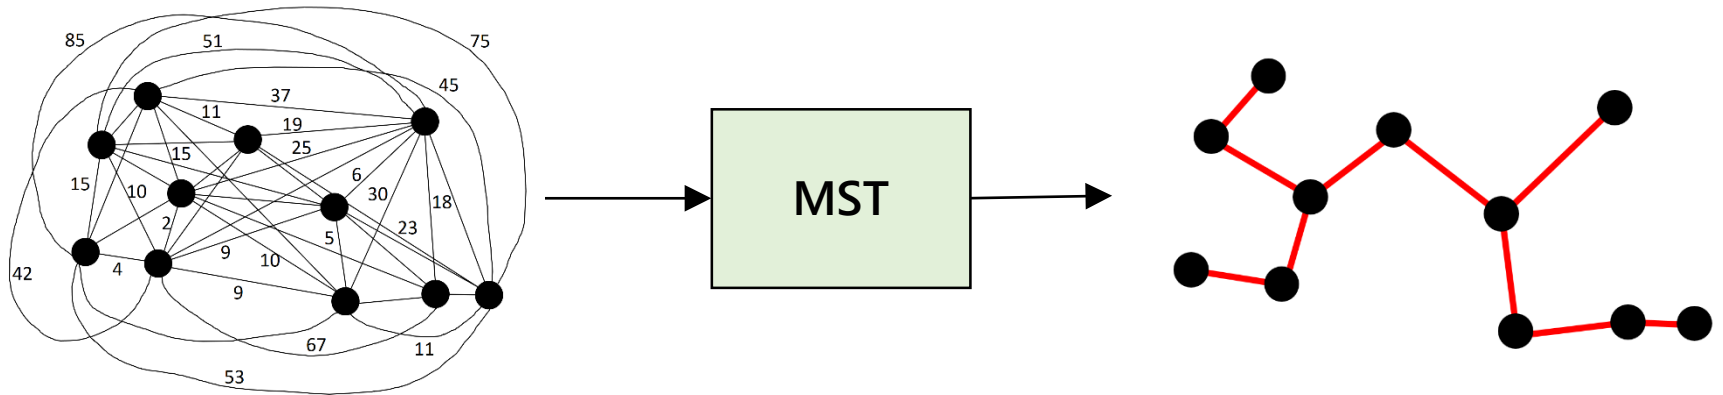
# Example: our little village

1. Represent the problem “cleverly” as a graph





# Example: our little village

## 2. Feed the graph to a Graph Algorithm



# Example: our little village

3. Use the output to determine the answer to your problem



## Solution:

Connect house A to B

Connect house B to C

Connect house C to D

Connect house D to Central

...

etc.





# Whew.

- Now we still need one of these:

**MST**

- In fact, we're going to look at two of them:

**Prim**

**Kruskal**

# Greedy Algorithms: Prim's Algorithm


Textbook: Chapter 9.1

# Idea of Prim's algorithm

- Initialize with any single vertex
- Each step: add one edge + one other vertex
- Definition of “best+feasible” greedy choice:
  - Smallest-weight edge
  - One vertex in the tree, one vertex NOT in the tree
  - → it does not create a cycle
- N-1 steps
- Result:
  - N vertices, N-1 edges, *connected* → spanning tree
  - This does not prove it is a *minimum* ST. (But it is.)


# Example 1

start  
vertex  
(root)




# Example 1

start  
vertex  
(root)




# Example 1

start  
vertex  
(root)




# Example 1

start  
vertex  
(root)



# Example 1


start  
vertex  
(root)






# Example 1

start  
vertex  
(root)




# Example 1

start  
vertex  
(root)




# Example 1

start  
vertex  
(root)




# Example 1

start  
vertex  
(root)




# Example 1

start  
vertex  
(root)



# Example 1 – completed MST

start  
vertex  
(root)



# Prim's algorithm

**Algorithm Prim(G)**

$V_T \leftarrow \{v_0\}$  // init tree with one (arbitrary) vertex

$E_T \leftarrow \emptyset$  // init tree with no edges

**for**  $i \leftarrow 1$  **to**  $|V|-1$  **do** // loop until all vertices added to tree


**find** a min-weight edge  $e=(u,v)$  **from**  $E$   
        **where**  $u$  **is in**  $V_T$  (in the tree)  
        **and**  $v$  **is in**  $V-V_T$  (not yet in the tree)

$V_T \leftarrow V_T \cup \{v\}$  // add the vertex  $v$  to the tree

$E_T \leftarrow E_T \cup \{e\}$  // add the edge  $(u,v)$  to the tree


**return**  $T = (V_T, E_T)$

# Example 2





## Example 2



# Greedy Algorithms: Kruskal's Algorithm


Textbook: Chapter 9.2

# Context

- Another one of several “greedy algorithms” we are examining:
  - Minimum Spanning Tree of a graph
    - Prim’s algorithm
    - Kruskal’s algorithm
  - Shortest Paths from a Single Source in a graph
    - Dijkstra’s algorithm
  - Graph coloring


# Kruskal's (overview)

- Repeatedly add a minimum-weight edge that does not introduce a cycle
- Example:




# Kruskal's (overview)

- Result (minimum spanning tree):




# Kruskal's algorithm (basic idea)

Kruskal( $G$ )



```
    sort edges of  $E$  in ascending order by weight
     $V_T \leftarrow V$                                 //  $T$  has all the vertices of  $G$ 
     $E_T \leftarrow \emptyset$                         // start with no edges in  $T$ 
    count  $\leftarrow 0$ 
     $k \leftarrow 0$                                 // index over edges of  $G$ 
    while count  $< |V|-1$  do                        // done when  $T$  has this many edges
         $k \leftarrow k + 1$ 
        if  $E_T \cup \{e_k\}$  is acyclic             // safe to add this edge to  $T$ ?
             $E_T \leftarrow E_T \cup \{e_k\}$         // ...then add it
            count  $\leftarrow$  count + 1
    return  $T = (V_T, E_T)$ 
```



These two bits are “efficiency challenges”

# Kruskal's algorithm: Implementation challenges

## 1. Sort the edges

- We know several  $O(N \log N)$  methods
- Which will serve us well?

## 2. Determine if adding an edge would create a cycle

- Maybe use a DFS or BFS to test for a cycle?
  - These are  $O(N^2)$  and we have to do it  $O(N)$  times
  - Can we improve on  $O(N^3)$ ?
- The answer is Yes, with a clever data structure

# Disjoint Subsets (aka “Union-Find”)

- A collection of disjoint subsets – any element can only be in one subset at any time
- Operations on a DS:
  - **Makeset(x)** – creates a new subset with the element x
  - **Find(x)** – returns the subset that contains x
  - **Union(x,y)** – merges the subsets containing x and y together



# DS/Union-Find Example

```
for each element x in {1,2,3,4,5,6,7,8}
    makeset(x)
```

→ DS is now {1} {2} {3} {4} {5} {6} {7} {8}

```
union(2,7)
```

→ DS is now {1} {2,7} {3} {4} {5} {6} {8}

```
union(1,4)
```

→ DS is now {1,4} {2,7} {3} {5} {6} {8}

```
y ← find(4)
```

→ y is now {1,4}

```
union(y,3)
```

→ DS is now {1,4,3} {2,7} {5} {6} {8}

```
x ← find(1)
```

→ x is now {1,4,3}

```
y ← find(7)
```

→ y is now {2,7}


```
union(x,y)
```

→ DS is now {1,4,3,2,7} {5} {6} {8}

# Kruskal's with disjoint subsets

- Maintain DS of vertices in the spanning tree  $T$
- Initially each vertex is a separate subset
- When an edge  $(u,v)$  is added to  $T$ :
  - $\text{DS.union}(u,v)$
- Each subset is a connected component
  - It's also a tree – a subset of the eventual MST
- If  $u,v$  are in the same subset *do not add edge*
  - It would create a cycle
- At the end there will be only one subset in DS
  - $T$  is a single connected component


## Another Kruskal example (using disjoint subsets)



- After the initialization
- PQ contains sorted list of edges
- DS has one subset for each vertex

PQ	Subsets						Solution
<u>key: value</u>	{a}	{b}	{c}	{d}	{e}	{f}	
3:ef							(a)      (b)      (c)
4:af							
4:cd							
5:ae							
6:bc							(f)      (e)      (d)
8:be							
9:bd							
10:ab							
12:de							
15:ce							


## Another Kruskal example (using disjoint subsets)



- After iteration 1
- edge ef has been added
- e, f subsets merged

PQ	Subsets						Solution
<u>key: value</u>	{a}	{b}	{c}	{d}	{e}	{f}	
3: ef	{a}	{b}	{c}	{d}	{e, f}		(a)                  (b)                  (c)
4: af							
4: cd							
5: ae							
6: bc							
8: be							
9: bd							
10: ab							
12: de							
15: ce							


## Another Kruskal example (using disjoint subsets)




- After iteration 2
- edge af has been added
- a, f subsets merged

PQ	Subsets					
<u>key: value</u>	{a}	{b}	{c}	{d}	{e}	{f}
3:ef	{a}	{b}	{c}	{d}	{e,f}	
4:af	{a,e,f}	{b}	{c}	{d}		
4:cd						
5:ae						
6:bc						
8:be						
9:bd						
10:ab						
12:de						
15:ce						

Solution




## Another Kruskal example (using disjoint subsets)




- After iteration 3
- edge cd has been added
- c, d subsets merged

PQ	Subsets					
<u>key: value</u>	{a}	{b}	{c}	{d}	{e}	{f}
3:ef	{a}	{b}	{c}	{d}	{e,f}	
4:af	{a,e,f}	{b}	{c}	{d}		
4:cd	{a,e,f}	{b}	{c,d}			
5:ae						
6:bc						
8:be						
9:bd						
10:ab						
12:de						
15:ce						

Solution




## Another Kruskal example (using disjoint subsets)




- *No change in iteration 4*
- *a and e are in the same subset*
- *edge ae is not added because it would cause a cycle*

PQ	Subsets					
<u>key: value</u>	{a}	{b}	{c}	{d}	{e}	{f}
3:ef	{a}	{b}	{c}	{d}	{e,f}	
4:af	{a,e,f}	{b}	{c}	{d}		
4:cd	{a,e,f}	{b}	{c,d}			
5:ae	<b>{a,e,f}</b>	{b}	{c,d}			
6:bc						
8:be						
9:bd						
10:ab						
12:de						
15:ce						




## Another Kruskal example (using disjoint subsets)



- After iteration 5
- edge bc has been added
- b, c subsets merged


PQ	Subsets					
<u>key: value</u>	{a}	{b}	{c}	{d}	{e}	{f}
3:ef	{a}	{b}	{c}	{d}	{e,f}	
4:af	{a,e,f}	{b}	{c}	{d}		
4:cd	{a,e,f}	{b}	{c,d}			
5:ae	{a,e,f}	{b}	{c,d}			
6:bc	{a,e,f}	{b,c,d}				
8:be						
9:bd						
10:ab						
12:de						
15:ce						

Solution






## Another Kruskal example (using disjoint subsets)



- After iteration 6
- edge be has been added
- N-1 edges added, main loop ends
- algorithm returns solution

PQ	Subsets
<u>key: value</u>	<u>{a}</u> <u>{b}</u> <u>{c}</u> <u>{d}</u> <u>{e}</u> <u>{f}</u>
3:ef	{a} {b} {c} {d} {e,f}
4:af	{a,e,f} {b} {c} {d}
4:cd	{a,e,f} {b} {c,d}
5:ae	{a,e,f} {b} {c,d}
6:bc	{a,e,f} {b,c,d}
8:be	{a,e,f,b,c,d}
9:bd	
10:ab	
12:de	
15:ce	

Solution



# Kruskal's algorithm with PQ + disjoint subsets

Algorithm Kruskal(G)

```
Add all vertices in G to T           // add v's but don't add e's
Create a priority queue PQ           // will hold candidate edges
Create a collection DS               // disjoint subsets
for each vertex v in G do
    DS.makeset(v)
for each edge e in G do
    PQ.add(e.weight, e )             // PQ of edges by min weight
while T has fewer than n-1 edges do
    (u,v) ← PQ.removeMin()           // get next smallest edge
    cu ← DS.find(u)
    cv ← DS.find(v)
    if cu ≠ cv then                  // be sure u,v are not in
        T.addEdge(u,v)               // the same subset
        DS.union(cu, cv)
return T
```

# Efficiency of Kruskal's

- With an efficient union-find algorithm, the slowest thing is the initial sort on edge weights
  - $O(|E| \log(|E|))$
  - Remember that  $|E|$  is (in the worst case)  $|V|^2$
  - So this is also  $O(|V|^2 \log(|V|))$
  - Since we usually use  $N$  as the number of vertices in a graph, this is  $O(N^2 \log N)$

# Prim's and Kruskal's TL/DR

- Same problem: Minimum Spanning Tree (MST)
- Both are greedy algorithms
- Both add edges one at a time
  - Prim's greedy choice: smallest edge that extends the tree
  - Kruskal's: smallest edge that doesn't make a cycle

# Greedy Algorithms: Dijkstra's Algorithm

Textbook: Chapter 9.3


# Context

- Another one of several “greedy algorithms” we will examine:
  - Minimum Spanning Tree of a graph
    - Prim’s algorithm
    - Kruskal’s algorithm
  - Shortest Paths from a Single Source in a graph
    - Dijkstra’s algorithm
  - Graph coloring

# Problem:


## Single-source Shortest Paths

- Find the shortest path from a chosen vertex (the *source*) to every other vertex




# What about BFS?

- Simple/basic BFS *already solves this* for an unweighted graph:



- ... but not for weighted graphs. Consider the distance between a and d:



- Algorithm to find shortest paths in weighted graphs needs to consider the weight on the edge before including it in the solution*




# Idea of Dijkstra's algorithm

- Remember the best-known shortest distances for all vertices
  - Initially “infinity” for all
- Choose the nearest unprocessed vertex
  - Definition of “nearest” tbd
- Look at all of its neighbors
- Update their known shortest distances (“Relax”)
- Repeat

# Relaxation


- Dijkstra refers to “relaxing” a vertex
- Meaning: update the best known shortest path to  $v$




We are at an intermediate stage:  
So far we “know” that we  
can get from  $s$  to  $u$  with cost 5  
and from  $s$  to  $v$  with cost 9

Using the new information about  
edge  $(u,v)$  we now know there is a  
cheaper path to  $v$

# Relaxation – another example



RELAX(*v*)




5+2 is no better than 6

No improvement,  
so no change this time


# Dijkstra Example

Find the shortest paths from A to all other vertices



# Dijkstra Example


**Initialize:**



$S: \{\}$

# Dijkstra Example

Add vertex A



# Dijkstra Example


Relax neighbors of A

Ⓐ

(dist)

*Q*:



<i>A</i>	<i>B</i>	<i>C</i>	<i>D</i>	<i>E</i>
0	$\infty$	$\infty$	$\infty$	$\infty$
	10	3	$\infty$	$\infty$



*S*: { *A* }

# Dijkstra Example

Add vertex C



(dist)



<i>Q:</i>	<i>A</i>	<i>B</i>	<i>C</i>	<i>D</i>	<i>E</i>
	0	$\infty$	$\infty$	$\infty$	$\infty$
		10	3	$\infty$	$\infty$

*S:* { *A*, *C* }



# Dijkstra Example

Relax neighbors of C





(dist)

Q:	A	B	C	D	E
	0	$\infty$	$\infty$	$\infty$	$\infty$
		10	3	$\infty$	$\infty$
		7		11	5

S: { A, C }

# Dijkstra Example

Add vertex E





(dist)

<i>Q:</i>	<i>A</i>	<i>B</i>	<i>C</i>	<i>D</i>	<i>E</i>
	0	$\infty$	$\infty$	$\infty$	$\infty$
		10	3	$\infty$	$\infty$
		7		11	5

$S: \{ A, C, E \}$

# Dijkstra Example

Relax neighbors of E





(dist)

<i>Q</i> :	<i>A</i>	<i>B</i>	<i>C</i>	<i>D</i>	<i>E</i>
	0	$\infty$	$\infty$	$\infty$	$\infty$
		10	3	$\infty$	$\infty$
		7		11	5
		7		11	

*S*: { *A*, *C*, *E* }

# Dijkstra Example

Add vertex B




(dist)

<i>Q:</i>	<i>A</i>	<i>B</i>	<i>C</i>	<i>D</i>	<i>E</i>
	0	∞	∞	∞	∞
		10	3	∞	∞
		7		11	5
		7		11	

*S:* { *A*, *C*, *E*, *B* }

# Dijkstra Example




Relax neighbors of B

(dist)


*Q:*

<i>A</i>	<i>B</i>	<i>C</i>	<i>D</i>	<i>E</i>
0	$\infty$	$\infty$	$\infty$	$\infty$
	10	3	$\infty$	$\infty$
	7		11	5
	7		11	
			9	




*S:* { *A*, *C*, *E*, *B* }

# Dijkstra Example



Add vertex D



(dist)

*Q:*

<i>A</i>	<i>B</i>	<i>C</i>	<i>D</i>	<i>E</i>
0	$\infty$	$\infty$	$\infty$	$\infty$
	10	3	$\infty$	$\infty$
	7		11	5
	7		11	
			9	

*S:* { *A*, *C*, *E*, *B*, *D* }

# Dijkstra's Algorithm

- Builds a tree of shortest paths rooted at the starting vertex
- This is a greedy algorithm: it adds the closest vertex, then the next closest, and so on (until all vertices have been added)

High-level pseudocode:

1. Initialise  $d$  and  $prev$
2. Add all vertices to a PQ with distance from source as the key
3. While there are still vertices in PQ
4.     Get next vertex  $u$  from the PQ
5.     For each vertex  $v$  adjacent to  $u$
6.         If  $v$  is still in PQ, relax  $v$

1.  $Relax(v)$  :
2.     if  $d[u] + w(u,v) < d[v]$
3.          $d[v] \leftarrow d[u] + w(u,v)$
4.          $prev[v] \leftarrow u$
5.         PQ.updateKey( $d[v]$ ,  $v$ )


# Output from Dijkstra's

- There are (at least) two possible outputs from Dijkstra's algorithm:
  - Tree of shortest paths from  $v$  to all other vertices
  - List (map) of total costs of shortest paths from  $v$  to all other vertices. I.e. the list tells you " $\text{min\_distance}(v, w)$ " for all the vertices reachable from  $v$ .



# Similarity of Dijkstra to Prim

- Both accumulate a tree  $T$  of edges from  $G$
- Each iteration: select the minimum priority edge adjacent to the tree that has been built so far
- In Prim's the priority of an edge is simply the weight of the edge



- In Dijkstra's the "priority" is the weight of the edge  $(u, v)$  plus the distance from the start to the parent of  $v$


# Sample application of Dijkstra's

- Suppose London wants fresh pineapples from Hawaii.
- There are no direct flights, but many possible connections.
- What is the best possible route to minimize overall shipping cost?

# Input: Shipping costs, city to city

- Honolulu to Chicago 105
- Honolulu to San Francisco 75
- Honolulu to Los Angeles 68
- Chicago to Boston 45
- Chicago to New York 56
- San Francisco to Boston 71
- San Francisco to New York 48
- San Francisco to Atlanta 63
- Los Angeles to New York 44
- Los Angeles to Atlanta 57
- Boston to London 88
- New York to London 65
- Atlanta to London 76




# Graph model of the problem



Apply Dijkstra's algorithm to find the cheapest cost from Hawaii to London  
(bonus: cheapest cost to all the other cities, too)


# Dijkstra limitation: negative weight edges

- Dijkstra's algorithm doesn't work with negative weight edges
- If we added a new edge to T, and it had a negative weight, then there could exist a shorter path (through this new vertex) to vertices already in T
- For example, consider graph A below.
  - Graph B is the result of running Dijkstra's algorithm on A.
  - But clearly there exists a path such as a-c-e in graph C that is shorter than the path found in B. Therefore Dijkstra's algorithm did not work on this graph that has a negative edge weight.



# Greedy Algorithms: Graph Coloring

Textbook: Mentioned several times, but not covered in-depth. Look in the index under “graph coloring”.



How many  
colors do you  
really NEED?

# Map coloring

- Problem: Color the regions on a map
  - Regions that share a border must be different colors
  - Meeting at a single point is not a border
- As a decision problem:
  - Can this map be colored with  $N$  colors?
- As an optimization problem:
  - What is the minimum number of colors needed to color this map?




# Graph representation

- One vertex for each region
- Edge between regions if they share a border
- Problem re-stated as a graph problem:
  - Assign colors to the vertices of a graph so that no adjacent vertices are the same color

# Graph coloring problem

- Color a graph with as few colors as possible such that no two adjacent vertices are the same color
- Example:





We say that this graph is *3-colorable*

# Graph coloring – greedy algorithm


- Start with just one color
- Consider the vertices in a specific order  $v_1, \dots, v_n$
- For each  $v_i$ , assign the first available color not used by any of  $v_i$ 's neighbours
- If all colors are in use by neighbours, add a new color

# Examples



# Is this algorithm optimal?

- Consider the previous graph but with vertices numbered differently



- Needed only two colors before
- The order of considering the vertices matters
- Greedy algorithms do not always yield optimal solutions
- But like brute-force, they are often worth considering because they may be easy to implement

# Puzzle – just for fun!

- Make a graph that represents a planar map and that *requires* 4 colors

# Practice problems

1. Chapter 9.1, page 324, question 9
2. Chapter 9.2, page 331, questions 1,2
3. Chapter 9.3, page 337, questions 1,2,4